

When Local Truth Produces Global Falsehood: Moral Grounding, Institutional Volition, and Systemic Error in Artificial Intelligence

Emanuel Shirbint

Independent Researcher, Tel Aviv, Israel

Corresponding author: shirbint@yahoo.com

ORCID: 0009-0004-8538-7140

Preprint. This manuscript has not undergone peer review.

Abstract

Evaluation of artificial intelligence concentrates on local performance — accuracy, calibration, robustness — while the upstream normative and volitional architecture that fixes what a system is for remains largely unexamined by those metrics. Existing scholarship has established that bias, proxy variables, problem formulation, measurement, alignment, and governance are value-laden; what remains insufficiently integrated is a single propagation account connecting moral grounding, authorized volition, teleological translation, representation, AI execution, and feedback. This article develops such an account: the Normative-Volitional Architecture of AI-Mediated Action, $M - W - T - R - E - I_{AI} - D - A - C$. Moral grounding (M) constrains what may legitimately be pursued; human or institutional volition (W) commits to a direction; teleological specification (T) translates that commitment into objectives and proxies; representation (R) determines which reality is available to the system; and AI performs inference and execution within that structure. The central argument is that AI executes an operational representation of authorized will; it does not independently legitimize ultimate ends. Systemic error therefore arises not only from malformed will but from the *normative translation gap* between defensible purposes and their operational encodings, and it persists under *normative closure* — the absence of an institutionalized feedback path through which consequences can reopen objectives, representations, and authority, rather than merely retraining models. Documented cases in healthcare allocation, risk assessment, hiring, engagement optimization, and welfare administration illustrate how locally correct outputs can constitute globally misleading trajectories, and a reflexive governance framework specifies remedies at each architectural layer.

Keywords: AI ethics; moral grounding; objective functions; systemic error; sociotechnical systems; normative feedback

1. Introduction: The Paradox of Correct Systems That Mislead

A provocative intuition circulates in public debate about artificial intelligence: that an amoral machine, however accurate, “always lies.” Whether any current AI behavior qualifies as lying in the intention-dependent philosophical sense remains contested, and the present argument does not depend on resolving that question. What matters is that the intuition indicates something real that the vocabulary of accuracy cannot capture: a system can be right about everything it measures and wrong about everything it is trusted for. This article reconstructs the intuition as

a defensible thesis about architecture rather than intention. A deployed AI system operates downstream of human and institutional acts of moral judgment, volitional commitment, problem formulation, representation, and authorization. When those upstream acts are morally malformed, procedurally illegitimate, or — the subtlest case — adequately motivated but defectively *translated* into operational objectives and representations, the system’s locally correct outputs can combine into a globally false, harmful, or misleading trajectory. AI does not execute morality directly. It executes an operational representation of human or institutional will; that will must itself be morally grounded, legitimately authorized, and adequately translated into objectives and representations.

The motivating paradox is empirical. A commercial algorithm widely used in United States health systems predicted future healthcare costs well; but because cost served as the proxy for medical need, and because less is historically spent on Black patients at equal levels of sickness, accurate cost prediction systematically deprioritized sicker but underserved patients (Obermeyer et al. 2019). No local metric registered the failure, because the failure was not in any inference. Nor, on the available evidence, was it in anyone’s malice: the plausible institutional intention — direct extra care to those who need it — was defensible. The failure occurred in the translation of that intention into a computable target and in the representation of patients through expenditure histories. Systemic error, this article argues, does not require an evil will. It can arise wherever the chain from moral grounding to machine action loses, replaces, or distorts meaning, and wherever consequences lack the institutional power to reopen the chain.

Three commitments frame the argument. First, no moral agency is attributed to machines; the article distinguishes moral agency from moral grounding, and instrumental subgoal formation from normative authorship (Sections 6 and 8). Advanced systems may plan, revise strategies, and generate intermediate goals; nothing in that capability confers legitimate authority to determine the ultimate ends of a sociotechnical system. Second, truth is not subordinated to morality: propositional truth values are independent of moral evaluation, and the concept of “global falsehood” is explicated strictly in terms of action-guiding adequacy, selection, and framing (Section 7). Third, the scope is bounded: the argument concerns socially consequential, action-guiding, institutionally embedded AI — systems allocating rights, resources, opportunities, risks, or burdens — not calculators or narrow physical-prediction models (Section 8.2). The stronger a system’s influence over persons, institutions, and material outcomes, the less adequate purely local measures of factual correctness become.

The article makes its contribution by integration rather than by any single novel observation. It binds together, in one architectural account: moral grounding; authorized volition; teleological translation; semantic and ontological representation; AI inference and execution; the distinction between local and systemic error; conditional capability-based amplification; the missing normative feedback loop; and the mechanism of systemic misleading. Two named concepts organize the integration: the *normative translation gap* — the distance between a morally articulated purpose and the formal objectives, proxies, representations, and decision rules through which it is operationalized — and *normative closure* — the condition of a system that can learn indefinitely how to pursue its objective without possessing an institutional mechanism for learning whether the objective should be pursued.

The article proceeds as follows. Section 2 positions the framework against existing scholarship and states the research gap. Section 3 distinguishes local from systemic error and states the scaling asymmetry in conditional form. Section 4 introduces the Normative-Volitional Architecture and its two feedback loops. Section 5 develops the four upstream levels — moral grounding, volition, teleological specification, representation — and the layered error taxonomy. Section 6 revises the concept of amorality. Section 7 analyzes local truth and global falsehood. Section 8 addresses AI's lack of independent normative authorship and the boundary conditions of the argument. Section 9 reanalyzes six cases through the architecture. Section 10 answers eight objections. Section 11 develops reflexive governance implications. Sections 12 and 13 discuss and conclude.

2. Research Gap and Theoretical Positioning

The claim that values enter technical systems is not new, and the framework developed here builds on several established traditions rather than discovering them. The philosophy of technology has long shown that artifacts settle political questions through design (Winner 1980) and that computer systems can embody preexisting, technical, and emergent bias (Friedman and Nissenbaum 1996), an insight elaborated programmatically by value-sensitive design (Friedman and Hendry 2019). Critical algorithm studies have documented harm across credit, policing, education, welfare, and search (O'Neil 2016; Eubanks 2018; Noble 2018). The fairness literature has moved upstream from output metrics to problem formulation — the choice of prediction target as a normative act (Passi and Barocas 2019) — and to measurement, showing that disputes coded as fairness disputes are often disputes about construct validity: whether the operationalized variable measures the concept it is taken to measure (Jacobs and Wallach 2021). Sociotechnical analysis has diagnosed the abstraction traps that arise when technical framing brackets the institutional context (Selbst et al. 2019). Alignment research has formalized objective misspecification and its hazards (Wiener 1960; Amodei et al. 2016; Russell 2019) and distinguished alignment to instructions, intentions, preferences, and values (Gabriel 2020). Organizational theory and cybernetics have long distinguished learning within goals from learning about goals: single- versus double-loop learning (Argyris and Schön 1978), the second-order observation of observing systems (von Foerster 2003), and reflexive governance of sociotechnical change (Voß, Bauknecht and Kemp 2006).

What remains insufficiently integrated is a single architectural account showing how moral grounding, institutional volition, teleological specification, representation, inference, action, and consequence form a continuous causal and normative chain — and how the absence of reflexive feedback allows locally correct systems to reproduce and amplify globally defective directions. The fairness literature analyzes T and R (targets, proxies, measurement) with precision but typically takes W as given and leaves M procedural. Alignment research analyzes the model side of T and the behavior of I_AI but tends to treat “human values” as an unanalyzed source rather than distinguishing grounding from volition from translation. Double-loop learning and reflexive governance supply the feedback distinction but were not developed for the specific propagation dynamics of machine-executed objectives, where capability scales the consequences of a fixed goal. This article's contribution is the integration: one propagation model in which each tradition's object of study occupies a definite location, and in which two concepts — the normative translation gap and normative closure — name the failure modes that appear only when the chain is viewed whole.

3. From Local Error to Systemic Error

Engineering practice recognizes a family of local failures: misclassification, miscalibration, hallucination, brittleness under distribution shift. Each is defined relative to a ground truth or specification taken as given; the apparatus of validation presupposes that the specification is sound. Wiener saw the limit of the presupposition at the origin of cybernetics: a machine pursues the purpose put into it, not the purpose meant (Wiener 1960). But the problem exceeds misspecification in the engineering sense, where a well-defined intended objective is imperfectly encoded. In socially consequential deployments the operational target is often chosen for availability, institutional convenience, or commercial measurability, inheriting values from datasets, incentives, and historical practice without any act of examined choice (Friedman and Nissenbaum 1996; Passi and Barocas 2019). This motivates a categorical distinction.

Local error. An incorrect fact, inference, classification, or prediction, identifiable by comparison with a ground truth internal to the system’s task definition, and often corrigible by more data, better models, and standard validation, within the limits set by irreducible task uncertainty.

Systemic error. A defective organizing principle — in the objective, the representation, the authorizing commitment, or the moral frame — that determines which goals, evidence, stakeholders, and outcomes the entire system treats as legitimate. Task-internal validation alone is insufficient to diagnose systemic errors, because the validation criteria are themselves products of the same organizing principle.

Table 1 contrasts the two classes.

Table 1. Local error versus systemic error in AI systems

Dimension	Local error	Systemic error
Location	Individual output (fact, inference, prediction)	Organizing principle: objective, representation, authorizing volition, moral frame
Reference standard	Ground truth internal to the task definition	Normative legitimacy of the task definition, representation, and authority
Detectability	Standard validation (test sets, benchmarks, calibration)	Task-internal metrics alone are insufficient; diagnosis requires normative, institutional, and stakeholder analysis
Effect of increased capability (objective, representation, and governance held fixed)	Tends to diminish	Execution of the error tends to gain scale, consistency, speed, and reach
Remedy	Retraining, debugging, better data	Reformulating purpose, target variable,

Dimension	Local error	Systemic error
Responsibility profile	Proximate source often traceable to model, data, implementation, or specification	representation, authorization, and governance Distributed across designers, institutions, deployers, and architecture
Feedback required	Operational (consequences → data → model)	Reflexive normative (consequences → representation, objective, volition, moral frame)
Example	Misclassifying one patient's risk	Treating predicted cost as a proxy for medical need across all patients

The scaling row requires careful statement, and the present formulation deliberately avoids two universal claims. It is not asserted that local error always decreases with capability, nor that systemic error always increases. The defensible claim is conditional: *holding the objective, representation, deployment authority, and governance structure fixed, greater capability can reduce local execution error while increasing the scale, consistency, speed, and reach with which a systemic error is implemented*. The condition matters because capability is not inherently blind to systemic error. More capable systems can, in principle, help detect proxy gaps, distributional harm, missing evidence, conflicting objectives, and long-horizon consequences. But this corrective role is realized only when the system — and, decisively, the institution deploying it — is authorized and designed to examine the objective rather than merely optimize it. The operative distinction is between *capability under a fixed objective* and *capability for reflexive objective revision*. The danger analyzed in this article arises when only the first increases: optimization power grows while the teleology, the represented world, and the authority structure remain immune to what the optimization reveals.

4. The Normative-Volitional Architecture of AI-Mediated Action

4.1 The model

The original insight that error can enter “at the objective function” compresses several distinct upstream acts into one symbol. An objective function does not appear from nowhere: someone, under some conception of what is permissible, commits an institution to a direction; that commitment is translated into targets and metrics; and the translation is executed over a constructed representation of the world. Unpacking the compression yields the architecture that organizes the remainder of the article:

M → W → T → R → E → I_AI → D → A → C, embedded throughout in **G**.

M — Moral grounding. The moral values, constraints, prohibitions, principles of justice, stakeholder obligations, and conditions of legitimacy under which an end may be pursued at all: human dignity, non-maleficence, non-domination, proportionality, responsibility, protection of vulnerable stakeholders, limits on acceptable means. M does not state what an actor wants; it determines whether a want is defensible.

Within this architecture, conscience is not an additional code of rules and, when the term is applied to AI, not a claim that the machine possesses an inner moral faculty. It names the orienting function of M: to practical intelligence what a compass is to a geologist or traveler, what the bubble in a spirit level is to a builder, and what a reference zero is to measurement. Rules and laws prescribe determinate conduct; conscience supplies the standpoint from which rules, exceptions, and combined consequences are judged. It is therefore not one norm among others but a second-order measure of the normative order itself—the point from which deviation, asymmetry, and accumulated imbalance become visible. In AI-mediated institutions, this function must be externalized and made contestable through M and G: explicit principles and prohibitions, stakeholder obligations, public justification, appeal, and reflexive revision. Conscience is not presumed infallible or reducible to private intuition; its judgments remain subject to public justification, contestation, and reflexive revision through M and G. Without such a reference point, intelligence may calculate direction and optimize movement with increasing precision while lacking any basis for determining whether the trajectory remains defensible. In the vocabulary of the architecture, the sequence is definite: conscience orients volition (M constrains and directs W); volition directs intelligence (W is translated through T and R); and intelligence selects means and executes action (I_AI, D, A). Direction is set before optimization begins, or it is not set at all.

W — Human or institutional volition. The authorized commitment to pursue a particular direction: selection of an end, acceptance of responsibility, authorization of action, and the standing decisions to continue, suspend, or revise. “Institutional volition” is used as shorthand for an authorized collective commitment embodied in governance decisions, laws, policies, procurement, management directives, mission statements, and allocations of authority and resources — not as the attribution of a mind to an organization. Crucially, volition is not assumed to be moral. It may be morally grounded; it may be malformed, explicitly immoral, procedurally illegitimate, unexamined, fragmented, imposed by market or bureaucratic incentive, or inconsistent with publicly declared values. Separating W from M is what makes the immoral or thoughtless will conceptually possible — and therefore analyzable.

T — Teleological specification. The translation of volition into operational form: formal objective functions, target variables, losses, rewards, performance metrics, KPIs, proxy variables, constraints, thresholds, time horizons, priority weights, and acceptance criteria. T is the site of a distinctive failure: an apparently moral intention can be *mistranslated*. The intention “identify medical need” can become the target “predicted expenditure” through nothing more sinister than data availability — and the resulting system is faithful to T while unfaithful to M and, possibly, to W.

R — Representation. The ontology and semantic model through which reality becomes available to the system: entities, categories, roles, relationships, events, states, risks, causal assumptions, features, labels, datasets, stakeholder models, temporal horizons, and the included and excluded forms of evidence. Representation is not a neutral mirror; it is a

constructed operational model of reality (a point long central to measurement critique: Jacobs and Wallach 2021). A system may pursue a legitimate objective within an impoverished representation that excludes vulnerable populations, historical injustice, unmet need, long-term harm, non-users, causal context, and the moral differences between superficially similar outcomes. What R omits, no downstream component can recover.

E — Evidence. The data, observations, documents, signals, labels, and measurements selected and weighted through R and T.

I_AI — AI inference, simulation, and optimization. Prediction, classification, reasoning, ranking, planning, generation, search, recommendation — including the construction of instrumental subgoals relative to the specified objective.

D — Decision. The selection or authorization of a course of action, with four distinguishable forms: machine-generated recommendation, automated decision, human-authorized decision, and institutionally authorized decision. Which form governs which stakes is itself a design and governance variable.

A — Action. Execution in the social, physical, informational, legal, medical, economic, or institutional world.

C — Consequences. Intended and unintended outcomes, stakeholder effects, distributions of benefit and burden, systemic and behavioral adaptation, and new social data.

G — Governance, responsibility, and contestability. Not a linear node but a cross-cutting layer intersecting every stage: authority, accountability, auditability, transparency, human oversight, participation, appeal, reversibility, responsibility assignment, regulatory constraint, and stakeholder representation. G determines whether the transitions between stages are documented, answerable, and revisable — or silent.

A one-sentence statement of the model's meaning: **AI executes neither human will in its fullness nor morality as such; it executes a formalized representation of an authorized objective.** An AI system does not receive morality. It receives data, prompts, objectives, policies, rewards, constraints, rules, examples, semantic structures, and authorization boundaries. It follows that even morally responsible humans can produce a morally defective system when the translation into machine-operable form is inadequate — the moral of the healthcare case, developed in Section 9.

4.2 Two feedback loops and normative closure

The architecture distinguishes two fundamentally different kinds of feedback, and the distinction carries much of the article's practical weight.

Operational feedback (C — E, C — I_AI, C — T at the level of thresholds and metrics). Consequences generate new data, retrain models, recalibrate thresholds, and improve performance. This is ordinary technical learning, and it asks: Did the model predict correctly? Did the intervention produce the expected result? Can the process become more accurate or efficient?

Reflexive normative feedback ($C \rightarrow R$, $C \rightarrow T$ at the level of target choice, $C \rightarrow W$, $C \rightarrow M$). Consequences trigger reconsideration of what the system represents, whether the objective remains legitimate, whether the original institutional commitment should continue, whether the moral assumptions were adequate, whether affected stakeholders were wrongly excluded, and whether responsibility and authority must be redistributed. This is not learning within a goal but learning about the goal, the distinction organizational theory captures as single- versus double-loop learning (Argyris and Schön 1978), second-order cybernetics as the observation of the observing system (von Foerster 2003), and governance theory as reflexivity (Voß, Bauknecht and Kemp 2006).

Define **normative closure** as the condition in which a system possesses strong operational feedback but lacks an institutionalized mechanism through which consequences can force reconsideration of the objective, the represented world, the legitimacy of the original will, or the moral framework itself. A normatively closed system can adapt indefinitely without learning whether it should continue to pursue the direction to which it is adapting. The claim defended here is deliberately non-universal: it is not asserted that consequences never reach M, but that *many deployed sociotechnical systems possess strong operational feedback while lacking an institutionalized, consequential, and contestable mechanism of normative feedback*. Appeals may exist formally but carry no authority; harms may be measured but translated only into technical bug reports; the model is retrained while the objective remains immune to revision. The contribution here is not the feedback distinction itself but its integration with objective-function analysis, representation, the local/systemic error distinction, capability scaling, and systemic misleading: normative closure is precisely the condition under which the conditional scaling claim of Section 3 becomes an operating reality. Two boundary points, developed in Section 8, already follow from the architecture itself (Fig. 1): AI may generate instrumental subgoals within it, and in currently deployed sociotechnical systems ultimate normative authority remains human and institutional.




Fig. 1 Normative-volitional architecture of AI-mediated action: moral grounding (M) – institutional volition (W) – teleological specification (T) – representation (R) – evidence (E) – AI inference and optimization (I_AI) – decision (D) – action (A) – consequences (C), intersected throughout by governance (G). Solid arrows: operational learning within the objective; dashed arrows: reflexive normative feedback that may reopen representation, objectives, authority, and moral grounding. Triangles mark systemic-error sites (normative malformation, translation failure, representational exclusion, missing reflexive loop). AI may generate instrumental subgoals; in currently deployed sociotechnical systems, ultimate normative authority remains human and institutional.

4.3 The Architecture of Meaning: how meaning enters — and leaks out of — AI systems

The model can be restated as an account of how meaning travels into computational systems, in four layers. A *normative-semantic layer* answers why the system exists, what is valuable, what future is pursued, which constraints are inviolable, who bears responsibility, and whose interests count (M and W). An *ontological-representational layer* answers what exists for the system: which entities, relations, events, states, risks, rules, and causal structures are recognized, which features and proxies stand for real-world concepts, and which realities remain invisible (T's semantic content and R). A *computational-operational layer* answers how the represented objective is pursued: processing, inference, prediction, simulation, optimization, decision support, automation, action (E, I_AI, D, A). A *consequence and revision layer* answers what actually happened, who benefited, who bore the cost, what was omitted, and whether the model, objective, representation, authority, or moral direction should be revised (C and the two loops).

Meaning does not enter AI systems magically. It is translated through a chain — moral judgment — volitional commitment — teleological specification — semantic and ontological representation — computational execution — material and social consequences — and at every transition it may be preserved, reduced, mistranslated, hidden, replaced by a proxy, subordinated to institutional convenience, distorted by incentives, or erased from the computational representation. Name the problem the **normative translation gap**: *the distance between a morally articulated purpose and the formal objectives, proxies, representations, and decision rules through which that purpose is operationalized*. The gap is not exotic; it is the ordinary condition of formalization, familiar from measurement theory (Jacobs and Wallach 2021) and from problem-formulation analysis (Passi and Barocas 2019). What the architectural view adds is the observation that the gap compounds across transitions — a small semantic substitution at T (need — cost) interacts with a representational restriction at R (patients — expenditure histories) — and that under normative closure the compounded gap is not merely undetected but actively reinforced, because operational feedback validates the system against its own translated targets.

The same point can be stated metrologically, and the analogy is offered strictly as an analogy. Every measurement system presupposes a reference that is not itself one of the measured quantities: the origin of a coordinate system, the zero of a scale, the level against which inclination is read. Conscience, as the orienting function of M, occupies this position in the architecture of practical intelligence: it is the zero point from which the adequacy of ends, translations, and trajectories is assessed, not an additional quantity computed within the system. Two consequences follow. First, the zero point cannot be derived from downstream operations: no refinement of T or R, and no amount of optimization at I_AI, generates the reference against which their results are judged, just as no set of measurements yields the origin of its own coordinate system. Second, a system that lacks the reference does not thereby become inaccurate; it becomes unanchored. Its outputs remain internally consistent — locally true in the sense of Section 7 — while the question of what the whole trajectory is measured from goes unasked. This is the metrological restatement of the article's central claim: local correctness is displacement measured within the system; global adequacy is position relative to a zero point the system does not itself supply. In Aristotelian terms, conscience so understood is

a condition of *practical truth* — the adequacy of action-guiding judgment to the ends it serves — and leaves propositional truth untouched (Aristotle 2009).

A note on philosophical provenance is owed. The framework is inspired in part by the Schopenhauerian priority of will over intellect — the thesis that willing establishes ends which the intellect serves rather than originates (Schopenhauer 1969) — while departing from Schopenhauer in a decisive respect: his metaphysical Will is blind striving, not morally responsible volition, and cannot be mapped onto organizational goal-setting or moral accountability. The architecture therefore takes from Schopenhauer only the structural priority of the volitional over the instrumental, and replaces blind striving with the tripartite analysis of morally accountable grounding (M), authorized commitment (W), and formal translation (T). Nothing in the argument depends on Schopenhauer's metaphysics.

5. Moral Grounding, Will, and Represented Objectives

5.1 Four levels that must not be conflated

The architecture's philosophical core is the refusal to conflate four levels that everyday and even scholarly usage runs together.

Moral grounding determines whether a desire, institutional goal, or proposed direction is morally defensible — under which conditions an end may be pursued, by which means, with which duties toward affected persons. *Volition* commits an actor or institution to a direction and assumes responsibility for it. *Teleological specification* renders the commitment computable. *Representation* constructs the world within which the computation operates. Each transition is a distinct act with a distinct failure mode, and the composite failure modes explain phenomena that single-level analyses cannot. If morality and will are fused — if every will is assumed to contain its own moral warrant — the immoral will becomes a conceptual impossibility, and the analysis loses exactly the cases (illegitimate procurement, mission–operation divergence, incentive-driven goal capture) it most needs to handle. If will and teleology are fused, the mistranslation of a defensible purpose into a defective target becomes invisible, and every systemic failure is misread as malice or negligence at the top. If teleology and representation are fused, the possibility of a legitimate objective pursued through an impoverished ontology disappears. Weber's separation of value rationality from instrumental rationality anticipates the first two distinctions (Weber 1978); Aristotle's separation of practical wisdom, which grasps ends, from cleverness, which merely reaches targets, anticipates the relation between M and everything downstream (Aristotle 2009). Measurement critique supplies the third distinction (Jacobs and Wallach 2021). The classical name for what anchors this hierarchy is conscience: it gives moral direction to volition, volition directs reason, and reason selects means and actions — a chain in which each level presupposes the one above it and none can substitute for the one above it.

Two consequences follow immediately. First, systemic error does not require an evil human intention: a morally defensible intention can coexist with teleological error, representational error, and local correctness. Second, an explicitly harmful will can coexist with flawless data, valid inference, and highly effective execution — in which case the AI system's technical quality is not a mitigating property but the mechanism of harm's efficiency.

5.2 A layered taxonomy of error

The architecture yields a taxonomy in which every error has an address (Table 2).

Table 2. Error locations in the Normative–Volitional Architecture

Error class	Location	Characteristic examples
Local computational error	I_AI, D, individual outputs	Incorrect prediction; hallucinated fact; misclassification; calibration failure
Evidential error	E	Incomplete data; sampling bias; measurement error; label error; stale evidence
Representational error	R	Excluded stakeholders; inadequate categories; missing causal context; reduction of complex human states to thin variables; unrepresented long-term consequences
Teleological error	T	Cost as proxy for need; arrests as proxy for offending; engagement as proxy for value; historical hiring as proxy for merit; efficiency defined without penalty for wrongful denial
Volitional / legitimacy error	W	Goal pursued without authorization from affected stakeholders; declared mission diverging from operational priority; incentives overriding declared commitments; responsibility diffused beyond answerability
Moral-grounding error	M	Persons treated merely as optimization variables; dignity and justice excluded from the objective space; stakeholders assigned negligible moral weight; no principled limits on means
Reflexive feedback error	$C \rightarrow R/T/W/M$	Harms measured but unable

Error class	Location	Characteristic examples
		to alter policy; appeals without authority; models retrained while objectives remain immune; complaints translated only into bug reports

Several errors typically coexist, and their interaction is the normal case rather than the exception: the healthcare system of Section 9.1 combines a defensible M and (plausibly) W with a teleological error at T, a representational restriction at R, and a reflexive feedback error that allowed the combination to persist until external researchers intervened. The taxonomy also disciplines the argument against overreach: not every imperfect objective is immoral, not every model error is a moral error, and the classes make it possible to say precisely which failure a given system exhibits instead of gesturing at “unethical AI.”

6. Amorality Revised: Value-Laden Is Not Morally Grounded

The predecessor concept of this article — amorality as the absence of moral criteria — requires refinement, because deployed AI systems are never value-free. They inherit values from target selection, measurement choices, institutional priorities, dataset construction, resource constraints, commercial incentives, law, historical practice, interface design, and authority structures (Winner 1980; Friedman and Nissenbaum 1996). The operative distinction is therefore not between value-laden and value-free systems but between value-laden and *morally grounded* ones. Define amorality accordingly: **the absence of explicit, justified, legitimate, and revisable moral grounding at the levels of objective formation, representation, authorization, and consequence evaluation.** An amoral system in this sense is typically saturated with values — implicit, inherited, unjustified, unexamined, concealed, procedurally illegitimate, morally incoherent, or reduced to institutional convenience.

Amorality so defined is one species of a broader genus, **normative malformation**, which the taxonomy of Section 5.2 articulates: moral absence, moral incoherence, hidden value inheritance, illegitimate volition, proxy distortion, representational exclusion, explicitly immoral objectives, and failure of normative feedback. The genus matters because remedies differ by species. Hidden value inheritance calls for disclosure and audit; proxy distortion calls for translation analysis at T; illegitimate volition calls for governance reform at W; explicit immorality calls for prohibition; feedback failure calls for institutional redesign of the reflexive loop. Collapsing them into a single accusation of “unethical AI” obscures both diagnosis and remedy.

7. Local Truth and Global Falsehood

7.1 The compositional problem

Truth attaches to propositions one at a time, but no one encounters a deployed system one proposition at a time. Users encounter rankings, scores, feeds, queues, and decisions: ensembles organized by T and R. The epistemically decisive question is not whether each element is true but whether the ensemble provides a defensible basis for action. The distinction can be stated

without any concession to relativism: *local truth concerns correspondence within a specified task; global adequacy concerns whether the task, representation, objective, and induced trajectory together provide a defensible basis for action*. Equivalently: a system may be locally correct relative to its labels and globally misleading relative to the social reality and moral purpose for which it is trusted. A fact remains true regardless of moral preference; but selection, ranking, framing, recommending, and acting are not truth-apt operations — they are normative operations performed continuously by every deployed system, and their adequacy is assessed against purposes, not labels.

The recurrent mechanisms by which correct outputs compose into misleading wholes map onto the architecture: selective framing and omission of morally relevant context (R); inappropriate proxies and defective success criteria (T); unequal stakeholder weighting (W and T); exclusion of long-term and diffuse consequences (R and T's horizons; Jonas 1984); normalization of historical injustice, whereby predictive learning converts descriptive regularity into prescriptive default (Fricker 2007; Noble 2018); and the transformation of persons into optimization variables, which removes the representational space in which dignity, voice, and contestation could register (Kant 1998). Hence: private truths can accumulate into a public falsehood when the architecture selecting and applying them is normatively defective — the falsehood residing not in any proposition but in the induced representation and trajectory.

7.2 Systemic misleading and borrowed authority

Refine the article's central term: **systemic misleading is a condition in which locally credible outputs lend epistemic and institutional authority to a defective objective, representation, or allocation of responsibility**. No intention is required; misleadingness is a functional property of outputs-in-context. The mechanism — call it *borrowed authority* — has five steps. (1) The model is visibly accurate in testable local domains. (2) Users and institutions generalize that credibility to the system's hidden normative assumptions. (3) The accuracy of execution legitimizes the objective: because the answers are right, the arrangement that selects the questions escapes scrutiny. (4) The objective's political and moral character is thereby obscured — it appears as a technical given rather than a contestable commitment. (5) Correctness functions as a shield against examination of purpose. Borrowed authority helps explain why systemic misleading can be especially difficult to detect and contest, even in the absence of intentional deception: a liar can be confronted and an interest exposed, whereas a systemically misleading arrangement recruits the system's genuine epistemic virtues into the concealment of its teleological and volitional defects. It also explains why the phenomenon intensifies under normative closure: operational feedback continuously renews the local credibility on which the borrowed authority draws, while nothing renews scrutiny of what the credibility is being lent to.

8. AI Without Independent Normative Authorship

8.1 Instrumental autonomy versus normative authorship

The claim that AI “executes represented will” must not be mistaken for the claim that AI has no agency of any kind. Contemporary agentic systems plan, select strategies, revise intermediate plans, interact autonomously with environments, and pursue goals over extended action sequences; they can generate intermediate plans and functionally instrumental subgoals

relative to specified objectives. The required distinction is between **instrumental goal formation** — constructing intermediate goals and strategies relative to a previously established objective — and **normative authorship** — the legitimate authority to determine which ultimate ends ought to govern a sociotechnical system and to justify those ends to affected persons. The first is a capability; the second is a normative status. Capabilities do not confer statuses: AI may possess functional autonomy in selecting means and instrumental subgoals without possessing legitimate normative authority over ultimate ends. Indeed, growth in instrumental autonomy *raises* rather than settles the authorship question, because more of the pathway from objective to consequence passes through machine-generated intermediate structure that no human deliberation has examined — which is why Section 11 places analysis of instrumental subgoals among the governance requirements at I_AI.

Two boundary clarifications prevent overreach. The argument does not assert that future AI could never possess any form of moral agency; that speculative question need not be resolved here. The present argument concerns current and near-term deployed AI systems whose authority, objectives, deployment domains, and normative standing remain institutionally derived. And the argument does not deny that AI can *contribute* to moral reflection — exposing contradictions, surfacing evidence of harm, simulating consequences (Objection 8, Section 10). Advisory capability for moral reflection is distinct from final authority over binding ends.

8.2 Scope: a graded boundary condition

Nor does the argument claim that every computational system requires a moral architecture in order to be reliable. A calculator is accurate without representing human dignity; a narrow weather model is epistemically reliable within its task. The thesis concerns socially consequential AI: action-guiding systems, decision support affecting persons, systems allocating rights, resources, opportunities, risks, or burdens, systems embedded in institutions, systems whose outputs shape representations of people and society, and systems capable of acting or authorizing action at scale. The boundary is graded, not binary: **the stronger the system's influence over persons, institutions, material resources, and socially consequential action, the less adequate purely local measures of factual correctness become.** Moral grounding requirements rise with stakes, scale, autonomy, irreversibility, opacity, vulnerability of affected populations, duration of effects, difficulty of contestation, and distance between operator and affected person. At the low end of every gradient, ordinary engineering validation suffices; at the high end, it captures only a limited subset of what matters for legitimate and responsible deployment.

9. Case-Based Illustrations

Each case is analyzed through the architecture: M (governing moral principles), W (authorizing commitment), T (operational translation), R (representation and exclusions), E/I_AI (evidence and computation), D/A (decision and action), C (consequences), and the feedback question — could consequences trigger only retraining, or also revision of purpose, representation, and commitment? Documented cases are cited; hypotheticals are marked. No accusation is made beyond the cited documentation.

9.1 Healthcare allocation (documented)

Obermeyer et al. (2019) analyzed a commercial risk-management algorithm used across US health systems. **M**: allocate scarce care-management resources according to medical need, with equal concern for equally sick patients. **W**: health systems' authorized commitment to identify patients for additional intervention — on its face defensible. **T**: the commitment was translated into predicting future healthcare cost, the available and monetizable variable; need was thereby operationalized as expenditure. **R**: patients became visible as expenditure histories; unmet need — sickness that generates little spending because access is blocked — had no representation. **E/I_AI**: billing data; accurate cost prediction (the study did not find the cost predictions inaccurate). **D/A**: enrollment prioritization for care programs. **C**: at equal risk scores, Black patients were substantially sicker; the authors estimated that correcting the target would roughly double the fraction of Black patients identified for additional help. **Feedback**: operational feedback was intact (the model predicted cost well and could be retrained); reflexive feedback was externally supplied — it took independent researchers with unusual data access to reopen T. The case is the canonical normative translation gap: moral concern for need, mistranslated as economic expenditure, executed accurately, under normative closure.

9.2 Recidivism risk assessment (documented, contested)

The COMPAS controversy (Angwin et al. 2016) and its scholarly aftermath (Chouldechova 2017; Kleinberg, Mullainathan and Raghavan 2017) illustrate normative conflict at every upstream level. **M**: liberty, public safety, proportionality, and equal treatment — values in genuine tension. **W**: jurisdictions' authority to define acceptable risk; the commitment to pretrial risk scoring was institutionally authorized but rarely deliberated with affected communities. **T**: offending was operationalized as rearrest — a proxy gap, since arrests reflect policing intensity as well as behavior; and among mathematically incompatible fairness criteria (calibration versus equalized error rates, given unequal base rates), a de facto selection was made by default rather than by public choice. **R**: defendants represented through criminal-history features; the differential enforcement history that generated those features was not represented. **C**: disparate error burdens across groups, each side of the controversy locally correct by its own metric. **Feedback**: the impossibility results show that no technical adjustment can dissolve the conflict; only reflexive feedback — a public, contestable choice among criteria and a re-examination of the rearrest proxy — addresses the actual error location, which is at T and W, not at I_AI.

9.3 Hiring (documented instance)

Reuters reported that Amazon discontinued an experimental résumé-screening model that downgraded indicators of being female, having learned from a decade of male-dominated outcomes (Dastin 2018). **M**: evaluate candidates on merit with equal opportunity. **W**: the declared goal — identify job suitability at scale. **T**: suitability was operationalized as resemblance to previously hired candidates; the past was silently equated with merit. **R**: the ideal employee was represented by historical exemplars; historical discrimination and unequal opportunity had no representation, so the model could not distinguish “was hired” from “merited hiring.” **C**: accurate imitation of the past functioning as systematic falsification of merit. **Feedback**: what is documented is that the problem was identified and development

discontinued before full deployment. On the present framework, the withdrawal can be interpreted as a pre-deployment reflexive intervention — a reopening of T and W rather than a retraining fix. So interpreted, the case illustrates both the failure mode and the remedy, which is why it is documented rather than accusatory.

9.4 Engagement optimization (documented pattern; effects partly contested)

M: serve users' informational interests and respect their autonomy; sustain rather than degrade the shared epistemic environment. **W:** platforms' commitment to maximize attention, retention, and revenue — a volition shaped by market incentive and not necessarily congruent with declared missions. **T:** user value operationalized as clicks, watch time, session length. **R:** the user represented as an engagement-generating profile; non-users, communities, and democratic institutions absent; long-term autonomy and epistemic quality outside the temporal horizon (Vallor 2016; Zuboff 2019). **I_AI:** among the best-solved prediction problems in applied machine learning. **C:** feeds in which arousing or polarizing content can dominate without any prediction being wrong. **Feedback:** operational feedback is exceptionally strong — the loop $C \rightarrow E \rightarrow I_AI$ runs continuously and strengthens the system — while normative feedback remains weak: the magnitude of societal effects is empirically contested, but the architectural point is independent of that controversy, since the objective's indifference to the difference between attention captured and value delivered is a design fact at T, and no internal mechanism exists through which evidence of harm reopens it.

9.5 Welfare allocation (documented pattern)

Eubanks (2018) documents automated eligibility and fraud-risk systems in US public services. **M:** subsistence rights, non-maleficence toward the vulnerable, procedural fairness. **W:** agencies' commitment — often legislatively driven — to reduce improper payments and caseload cost; the operative institutional will hidden behind ostensibly technical eligibility rules. **T:** asymmetric objectives — strong penalty for improper payment, weak or no penalty for wrongful denial — so the system's errors are structurally aimed at claimants. **R:** claimants represented primarily as fraud risk and cost; need, dependency, and the consequences of erroneous denial thinly represented or absent. **C:** burdens concentrated on the least powerful stakeholders. **Feedback:** contestability formally exists but is often ineffective for affected populations; complaints are absorbed as case-level exceptions rather than as evidence against the objective — a textbook reflexive feedback error.

9.6 High-stakes autonomous action (clearly hypothetical)

Consider, hypothetically, a military target-recognition system optimized for identification accuracy and engagement speed under encoded rules of engagement. **M:** proportionality and discrimination in the *jus in bello* sense; accountability for error. **W:** command authority — the one domain where the chain of authorization is doctrinally explicit, which is precisely what makes the case analytically clean. **T:** legitimacy operationalized as confidence thresholds and latency; the moral judgment of proportionality cannot be reduced without remainder to confidence thresholds, latency, or other computable proxies. **R:** the represented battlefield — sensor categories, combatant/non-combatant classifications — determines what can be discriminated at all. **D/A:** the decisive design variable is the separation between prediction and authorization to act; compressing decision tempo below the speed of human deliberation does

not remove moral judgment but silently replaces it with threshold arithmetic. **C:** irreversible. **Feedback:** after-action review exists institutionally, but under tempo pressure the reflexive loop is the first casualty. The case marks the limit thesis: as stakes and irreversibility rise, the gap between local correctness and global legitimacy widens, and the graded boundary condition of Section 8.2 reaches its maximum.

Across the six cases the invariant is now more precise than “the objective was wrong.” In one case (9.3) a pre-deployment reflexive intervention occurred; in one (9.1) reflexive scrutiny had to be imported from outside; in the others it is weak, contested, or structurally absent. The decisive variable is not the presence of error — every complex system errs — but whether the architecture permits consequences to reopen commitments.

10. Objections and Replies

Objection 1: AI generates goals and acts autonomously, so “AI merely executes human will” is false. Reply. The formulation to defend is not “merely executes” but “executes a formalized representation of an authorized objective,” which is compatible with rich instrumental autonomy. The objection conflates instrumental goal formation with normative authorship (Section 8.1). Systems that plan and generate subgoals do so relative to established objectives; the capability, however sophisticated, does not confer legitimate authority over ultimate ends, and its growth strengthens rather than weakens the case for explicit upstream architecture, since more consequential structure becomes machine-generated and unexamined.

Objection 2: Institutions do not literally possess will; “institutional volition” is a category error. Reply. Agreed, if taken metaphysically — which is why Section 4.1 defines institutional volition non-mentalistically, as authorized collective commitment embodied in governance decisions, laws, policies, procurement, directives, and allocations of authority and resources. The term names an observable institutional fact (that commitments are made, resourced, and enforced), not a group mind. Legal systems have long managed the same construction for corporate intention without metaphysical embarrassment.

Objection 3: Human will is frequently immoral, irrational, divided, or manipulated — so grounding AI in human will grounds it in sand. Reply. The objection supports the framework rather than refuting it. It is precisely because volition may be malformed that M and W must be distinguished: the architecture does not treat will as self-legitimizing but as requiring moral and institutional justification. A framework that fused morality and will could not even state this objection; the present framework makes it a named error class (volitional/legitimacy error, Table 2) with named remedies (Section 11, at W).

Objection 4: The model merely moves the alignment problem one level up — who aligns the humans and institutions? Reply. Partly accepted, and deliberately so. The framework does not promise an algorithmic solution to morality, and it would be suspect if it did. Its claim is locational: it identifies the loci — M, W, T, R, and the reflexive loop — where normative responsibility must remain explicit, contestable, and institutionally revisable, instead of being silently absorbed into technical artifacts where it becomes invisible and unrevisable. Keeping the problem visible at the level where politics, law, and ethics can engage it is not a failure to solve it; it is the precondition of addressing it by legitimate means.

Objection 5: Moral pluralism prevents any single legitimate objective. Reply. The framework requires no comprehensive doctrine. Abstention does not honor pluralism, since an unexamined objective is value-laden by inheritance (Section 6). What is required is thinner: minimum constraints already institutionalized in human rights law and professional ethics (dignity, non-maleficence, non-domination, proportionality, accountability, contestability); procedural legitimacy — explicit, publicly justified, contestable choices where reasonable disagreement persists, as among incompatible fairness criteria (Habermas 1996; Floridi et al. 2018); stakeholder participation; and, where appropriate, plural objective structures rather than a single scalar target.

Objection 6: AI might eventually become a moral agent, dissolving the human/machine authority distinction. Reply. Possibly; the argument does not depend on permanently denying it. The claims here address deployed systems whose normative authority remains derived from human institutions — which covers current and near-term socially consequential AI. If genuinely morally accountable artificial agents ever exist, the architecture would require extension, not the present analysis retraction: the questions of grounding, authorization, translation, and reflexive revision would then arise for those agents too.

Objection 7: Narrow technical systems do not need morality to be accurate. Reply. Accepted, and built into the framework as the graded boundary condition of Section 8.2. The thesis concerns socially consequential, action-guiding, institutionally embedded systems; the moral-architecture requirement scales with stakes, autonomy, irreversibility, and the vulnerability of those affected. The calculator objection defeats a universal claim the article does not make.

Objection 8: More capable AI could identify moral errors better than humans — so capability is the remedy, not the risk. Reply. The possibility is acknowledged and is already encoded in the conditional scaling claim (Section 3): capability can help detect proxy gaps, distributional harm, and long-horizon consequences — when the system and its institution are authorized and designed to examine the objective. The needed distinction is between advisory capability for moral reflection and legitimate authority to determine binding ends. AI can contribute evidence, expose contradictions, simulate consequences, and thereby strengthen the reflexive loop; it does not thereby acquire final normative authorship, for the same reason that an excellent expert witness does not thereby become the court.

11. Toward Reflexively Governed AI Architecture

If systemic error can enter at M, W, T, or R, and persist through normative closure, remedies must address each layer and, above all, institutionalize the reflexive loop. Post-hoc ethical review inspects outputs after the teleology has silently settled; the following requirements move governance upstream and keep it open. They extend the responsible-AI literature (Dignum 2019; Floridi et al. 2018; Friedman and Hendry 2019; Russell 2019) by assigning every requirement an architectural address.

At M — moral grounding. Explicit ethical purpose, published; minimum moral constraints encoded as upstream requirements rather than downstream filters; enumerated stakeholder obligations, including vulnerable and indirect stakeholders; prohibited ends and means; public justification of the system's existence.

At W — volition and authorization. Identification of who authorized the objective; documented institutional responsibility resistant to diffusion; accountable approval with named authority; specified conditions for suspension or termination; a record of disagreements and minority concerns.

At T — teleological specification. Target-variable justification as a standard artifact, on a par with test-performance documentation; proxy-gap analysis stating explicitly what is measured versus what is meant; comparison of alternative objectives; multi-objective design where a scalar target would force illegitimate trade-offs; explicit trade-off and time-horizon documentation; stress testing of success metrics; prohibition of proxies whose gap is known to encode injustice.

At R — representation. Ontology and category audit; verification that affected stakeholders — including non-users — are represented as parties, not only as variables; causal-context documentation; analysis of omitted variables and excluded realities; representation of long-term and indirect effects; explicit uncertainty about constructs; semantic traceability from moral purpose to operational variable.

At E — evidence. Source provenance; data quality and label validity; selection-bias analysis; documentation of the historical institutional processes that generated the data; explicit distinction between observed behavior and the morally relevant reality it imperfectly indicates.

At I_AI — execution. Uncertainty disclosure as a first-class output, limiting borrowed authority to validated domains; bounded optimization; monitoring for reward hacking; evaluation outside the narrow target metric; distribution-shift detection; analysis of machine-generated instrumental subgoals; constraints on autonomous action proportional to stakes.

At D and A — decision and action. Separation between prediction and authorization to act, strengthening with stakes; human review proportional to consequence; effective appeal and reversibility; traceable authority for each decision class; intervention thresholds; safeguards against automation bias.

At C — consequences. Distributional impact monitoring; long-term and non-user effects; behavioral and institutional adaptation; public reporting; stakeholder complaint mechanisms that classify complaints by error address (Table 2) rather than defaulting all of them to the technical layer.

For the reflexive loop — against normative closure. Consequences must be able to reopen the objective: review bodies with real authority over T, W, and deployment; periodic reconsideration of system purpose on a mandated schedule; sunset clauses requiring reauthorization rather than permitting drift; suspension powers; documented objective-revision protocols; stakeholder participation in reauthorization; and documented moral and institutional learning, distinct from model retraining, so that what was learned about the goal is preserved and not only what was learned within it.

The framework reframes the relation between ethics and engineering. An engineering practice that cannot state, justify, and revise its objective is incomplete as engineering: reliability is relative to purposes, and unexamined purposes leave reliability undefined. Moral grounding

and reflexive governance are to socially consequential intelligent systems what statics is to bridges — not external limitations on the discipline but part of its content.

12. Discussion

Contribution. The article’s contribution is integrative and locational. It unpacks the objective function into four upstream acts (grounding, volition, translation, representation), assigns every recognized failure mode an address in a single propagation architecture, and names the two failure conditions visible only at the architectural scale: the normative translation gap, which compounds across transitions, and normative closure, under which operational learning entrenches what reflexive learning should be questioning. Relative to problem-formulation and measurement scholarship, it adds the volitional level and the feedback dynamics; relative to alignment research, it decomposes the “human values” terminus into grounding, commitment, and translation; relative to double-loop learning and reflexive governance, it supplies the AI-specific propagation and scaling analysis. The account also preserves and sharpens the predecessor concepts: the local/systemic error distinction now carries a conditional rather than universal scaling claim; systemic misleading is explicated through the five-step mechanism of borrowed authority; and the missing C – M loop is restated in the defensible form that many deployed systems possess strong operational feedback while lacking institutionalized, consequential, and contestable normative feedback.

Limitations. Six should be acknowledged. First, the argument is conceptual; it presents no new empirical findings, and some illustrative claims — notably the societal effects of engagement optimization — remain empirically contested, though the architectural points do not depend on resolving them. Second, the model is semi-formal; formalizing the normative translation gap (for example, as a measurable divergence between construct and operationalization under a stated purpose) is an open task that may be only partially achievable. Third, the minimum constraints invoked against the pluralism objection are defended as overlapping consensus, not derived from first principles, and their margins will be contested; the framework deliberately routes that contest to legitimate procedure rather than resolving it by stipulation. Fourth, the governance requirements of Section 11 presuppose institutional capacity and impose real costs; the framework identifies what reflexive governance requires, not a costless path to it, and the political economy of who bears those costs is beyond its scope. Fifth, the analysis addresses deployed systems with identifiable objectives; foundation models complicate the picture with nested objectives (pre-training, fine-tuning, system prompts, downstream use) and divided custody of M, W, and T across developer, deployer, and user — an extension deferred to future work. Sixth, the architecture is stated as if the stages were cleanly separable; in practice M, W, T, and R interpenetrate, and the model’s linearity is an analytical idealization whose value lies in making failure addresses nameable, not in describing organizational process realistically.

Implications. For AI ethics, the account redirects the primary audit question from “is it accurate?” and even “is it fair?” to “what is it for, who committed to it, how was the commitment translated and represented, and can consequences reopen it?” For technical practice, target-variable justification, proxy-gap analysis, and semantic traceability should become standard documentation. For regulation, the analysis supports impact-assessment and reauthorization regimes that reach problem formulation, procurement, and objective revision — the stages at which, on this account, the decisive errors are committed and entrenched.

13. Conclusion

The argument resolves into a division of labor that no increase in computational capability can erase. Morality does not calculate the path; it evaluates and constrains the direction. Volition does not establish moral legitimacy by itself; it commits an actor or institution to a direction for which responsibility must be assumed. Teleological specification translates that direction into operational objectives — and may distort it. Representation determines which reality becomes visible to the system at all. Artificial intelligence makes execution, prediction, and optimization more precise; precision cannot, by itself, establish that the direction is legitimate. A system may therefore be locally accurate while executing a globally defective moral and institutional trajectory, and it may generate instrumental subgoals of great sophistication without acquiring legitimate normative authorship over ultimate ends. If the trajectory is to be corrigible, consequences must be able to reopen not only data and models but goals, representations, authority, and moral assumptions. Reliable socially consequential AI accordingly requires more than alignment of model behavior: it requires reflexive governance of the entire chain from moral grounding to consequence.

Intelligence makes movement more precise. Moral grounding determines whether the destination is defensible. Human and institutional volition authorize the journey. Representation defines the world through which the system can move. Artificial intelligence executes the represented path. Conscience, as the zero point of this chain, is not one rule within the system but the reference against which the practical adequacy and moral defensibility of its ends and decisions are assessed. Responsibility therefore remains with those who establish, authorize, formalize, and preserve the direction. When moral orientation is absent, concealed, or mistranslated in the architecture of goals, artificial intelligence can transform local truth into an efficient instrument of systemic error.

Author Contributions

Conceptualization, theoretical framework, methodology, literature analysis, visualization, writing—original draft, and writing—review and editing: Emanuel Shirbint. The author has read and approved the final manuscript.

Declarations

Funding. The author received no funding for this work.

Competing Interests. The author declares no competing interests.

Ethics Approval and Consent. Not applicable. This is a conceptual article involving no human participants, personal data, animal research, or new empirical data collection.

Data Availability. No new datasets were created or analyzed in this conceptual study. All examples and arguments are based on publicly available literature cited in the manuscript.

Generative AI Use. Generative AI tools were used to assist with the conceptual organization, drafting, linguistic revision, and preparation of the text-based flow diagram. The author critically reviewed and substantively revised all AI-assisted material, verified the arguments,

terminology, factual claims, references, and visual elements retained in the final manuscript, and assumes full responsibility for the content.

References

- Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). Concrete problems in AI safety. *arXiv preprint* arXiv:1606.06565. <https://doi.org/10.48550/arXiv.1606.06565>
- Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016). Machine bias. *ProPublica*, 23 May 2016.
- Argyris, C., & Schön, D. A. (1978). *Organizational learning: A theory of action perspective*. Addison-Wesley.
- Aristotle. (2009). *The Nicomachean ethics* (D. Ross, Trans.; L. Brown, Ed.). Oxford University Press.
- Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2), 153–163. <https://doi.org/10.1089/big.2016.0047>
- Dastin, J. (2018). Amazon scraps secret AI recruiting tool that showed bias against women. *Reuters*, 10 October 2018.
- Dignum, V. (2019). *Responsible artificial intelligence: How to develop and use AI in a responsible way*. Springer.
- Eubanks, V. (2018). *Automating inequality: How high-tech tools profile, police, and punish the poor*. St. Martin's Press.
- Floridi, L., Cowls, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., Luetge, C., Madelin, R., Pagallo, U., Rossi, F., Schafer, B., Valcke, P., & Vayena, E. (2018). AI4People—An ethical framework for a good AI society: Opportunities, risks, principles, and recommendations. *Minds and Machines*, 28(4), 689–707. <https://doi.org/10.1007/s11023-018-9482-5>
- Fricker, M. (2007). *Epistemic injustice: Power and the ethics of knowing*. Oxford University Press.
- Friedman, B., & Hendry, D. G. (2019). *Value sensitive design: Shaping technology with moral imagination*. MIT Press.
- Friedman, B., & Nissenbaum, H. (1996). Bias in computer systems. *ACM Transactions on Information Systems*, 14(3), 330–347. <https://doi.org/10.1145/230538.230561>
- Gabriel, I. (2020). Artificial intelligence, values, and alignment. *Minds and Machines*, 30(3), 411–437. <https://doi.org/10.1007/s11023-020-09539-2>
- Habermas, J. (1996). *Between facts and norms: Contributions to a discourse theory of law and democracy* (W. Rehg, Trans.). MIT Press.
- Jacobs, A. Z., & Wallach, H. (2021). Measurement and fairness. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FACCT '21)*, 375–385. <https://doi.org/10.1145/3442188.3445901>

- Jonas, H. (1984). *The imperative of responsibility: In search of an ethics for the technological age*. University of Chicago Press.
- Kant, I. (1998). *Groundwork of the metaphysics of morals* (M. Gregor, Trans. & Ed.). Cambridge University Press.
- Kleinberg, J., Mullainathan, S., & Raghavan, M. (2017). Inherent trade-offs in the fair determination of risk scores. In *Proceedings of the 8th Innovations in Theoretical Computer Science Conference (ITCS 2017)*, LIPIcs 67, Article 43, 43:1–43:23. <https://doi.org/10.4230/LIPIcs.ITCS.2017.43>
- Nissenbaum, H. (1996). Accountability in a computerized society. *Science and Engineering Ethics*, 2(1), 25–42. <https://doi.org/10.1007/BF02639315>
- Noble, S. U. (2018). *Algorithms of oppression: How search engines reinforce racism*. New York University Press.
- Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447–453. <https://doi.org/10.1126/science.aax2342>
- O’Neil, C. (2016). *Weapons of math destruction: How big data increases inequality and threatens democracy*. Crown.
- Passi, S., & Barocas, S. (2019). Problem formulation and fairness. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19)*, 39–48. <https://doi.org/10.1145/3287560.3287567>
- Russell, S. (2019). *Human compatible: Artificial intelligence and the problem of control*. Viking.
- Schopenhauer, A. (1969). *The world as will and representation* (E. F. J. Payne, Trans., 2 vols.). Dover Publications.
- Selbst, A. D., boyd, d., Friedler, S. A., Venkatasubramanian, S., & Vertesi, J. (2019). Fairness and abstraction in sociotechnical systems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19)*, 59–68. <https://doi.org/10.1145/3287560.3287598>
- Vallor, S. (2016). *Technology and the virtues: A philosophical guide to a future worth wanting*. Oxford University Press.
- von Foerster, H. (2003). *Understanding understanding: Essays on cybernetics and cognition*. Springer.
- Voß, J.-P., Bauknecht, D., & Kemp, R. (Eds.). (2006). *Reflexive governance for sustainable development*. Edward Elgar.
- Weber, M. (1978). *Economy and society: An outline of interpretive sociology* (G. Roth & C. Wittich, Eds.). University of California Press.
- Wiener, N. (1960). Some moral and technical consequences of automation. *Science*, 131(3410), 1355–1358. <https://doi.org/10.1126/science.131.3410.1355>

Winner, L. (1980). Do artifacts have politics? *Daedalus*, 109(1), 121–136.

Zuboff, S. (2019). *The age of surveillance capitalism: The fight for a human future at the new frontier of power*. PublicAffairs.